

Pronóstico

Forecasts

```
library(tidyverse)
```

```
— Attaching core tidyverse packages — tidyverse 2.0.0 —
✓ dplyr      1.1.4      ✓ readr      2.1.5
✓ forcats    1.0.0      ✓ stringr    1.5.1
✓ ggplot2    3.5.2      ✓ tibble     3.3.0
✓ lubridate  1.9.4      ✓ tidyr      1.3.1
✓ purrr      1.1.0

— Conflicts — tidyverse_conflicts() —
✖ dplyr::filter() masks stats::filter()
✖ dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(fpp3)
```

```
Registered S3 method overwritten by 'tsibble':
  method          from
  as_tibble.grouped_df dplyr

— Attaching packages — fpp3 1.0.1 —
✓ tsibble      1.1.6      ✓ feasts      0.4.1
✓ tsibbledata  0.4.1      ✓ fable       0.4.1

— Conflicts — fpp3_conflicts —
✖ lubridate::date() masks base::date()
✖ dplyr::filter()   masks stats::filter()
✖ tsibble::intersect() masks base::intersect()
✖ tsibble::interval() masks lubridate::interval()
✖ dplyr::lag()       masks stats::lag()
✖ tsibble::setdiff() masks base::setdiff()
✖ tsibble::union()   masks base::union()
```

Let's review this time series for Australian production. You can see that the gas production doesn't have constant variance and needs a mathematical transformation to stabilize it.


Assuming we'll make a forecast for 3 years of Aus production, let's separate our data in train and test.

```
gas_train <- aus_production |>
  filter_index(. ~ "2007 Q2")
```

Now let's visualize the time series:

```
gas_train |>
```

```
autoplot(Gas)
```




Lets use box-cox with Guerrero feature and see the changes.

```
gas_lambda <- gas_train |>  
  features(Gas, features= guerrero) |>  
  pull()
```

```
gas_lambda
```

```
[1] 0.1037006
```

```
gas_train |>  
  autoplot(box_cox(Gas, gas_lambda))
```



¿Which forecast method is the most appropriate for this time series?

Note that all the mathematical transformations have to be done inside the fitting function, not before.

```
#tabla con los modelos ajustados
gas_fit <- gas_train |>
  model(
    drift = RW(box_cox(Gas, gas_lambda) ~ drift()),
    snaive = SNAIVE(box_cox(Gas, gas_lambda)),
    media = MEAN(box_cox(Gas, gas_lambda))
  )

gas_fit
```

```
# A mable: 1 x 3
      drift  snaive  media
  <model> <model> <model>
1 <RW w/ drift> <SNAIVE> <MEAN>
```

Lets make a diagnosis of the models residuals. The function `augment()`, it allows us to obtain the residuals and other adjusted values of the models.

```
gas_aug <- gas_fit |>
  augment()
```

gas_aug

```
# A tsibble: 618 x 6 [1Q]
# Key:           .model [3]
  .model Quarter  Gas .fitted .resid .innov
  <chr>    <qtr> <dbl>    <dbl> <dbl> <dbl>
1 drift  1956 Q1  5      NA     NA     NA
2 drift  1956 Q2  6.00    5.12  0.884  0.190
3 drift  1956 Q3  7.00    6.14  0.863  0.160
4 drift  1956 Q4  6.00    7.16 -1.16  -0.214
5 drift  1957 Q1  5      6.14 -1.14  -0.245
6 drift  1957 Q2  7.00    5.12  1.88   0.377
7 drift  1957 Q3  7.00    7.16 -0.157 -0.0272
8 drift  1957 Q4  6.00    7.16 -1.16  -0.214
9 drift  1958 Q1  5      6.14 -1.14  -0.245
10 drift 1958 Q2  7.00    5.12  1.88   0.377
# i 608 more rows
```

Residual Diagnosis

A good forecast model will produce residuals with the following characteristics:

1. **Residuals are not autocorrelated:** if correlations are detected between residuals, there's still useful information yet to be modeled.
2. **The residuals average is zero:** If the average is different than zero, then the forecast is biased.


Additional to this, the following characteristics are useful, not determinant:

3. The residuals have a constant variance
4. The residuals are *normally* distributed

Gráfica de los residuos de los modelos:

```
gas_aug |>
  autoplot(.innov) +
  facet_wrap(~.model, ncol=1, scales = "free_y")
```


Warning: Removed 5 rows containing missing values or values outside the scale range (`geom_line()`).



Gráfica edel ACF del residuo de los modelos:

```
gas_aug <- gas_fit |>
  augment()

gas_aug |>
  ACF(.innov) |>
  autoplot() +
  facet_wrap(~.model, ncol = 1, scale = "free_y")
```




Gráfica del histograma de los residuos

```
gas_aug |>
  ggplot(aes(x = .innov, fill = .model)) +
  geom_histogram() +
  facet_wrap(~.model, ncol=1, scales = "free_y")
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Warning: Removed 5 rows containing non-finite outside the scale range
(`stat_bin()`).



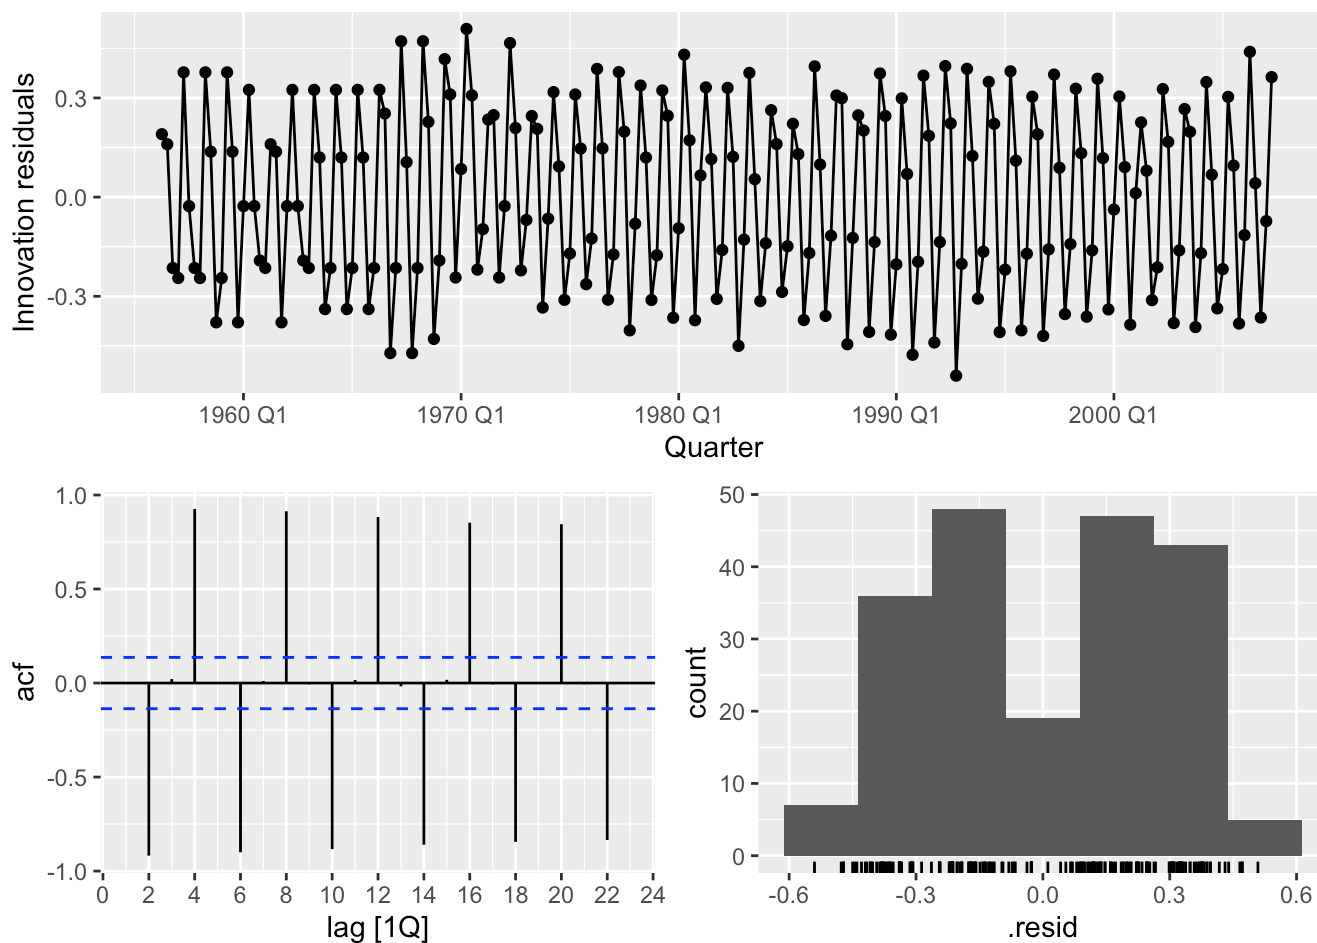
Otra opción

```
gas_fit |>
  select(drift) |>
  gg_tsresiduals()
```

Warning: Removed 1 row containing missing values or values outside the scale range (`geom_line()`).

Warning: Removed 1 row containing missing values or values outside the scale range (`geom_point()`).

Warning: Removed 1 row containing non-finite outside the scale range (`stat_bin()`).



```
aus_prod_recent <- aus_production |>
  filter_index("1999 Q1" ~ .)
```

```
gas_fc <- gas_fit |>
  forecast(h="3 years") # 12 Q es 3 años, se puede escribir solo 12

gas_fc
```


```
# A fable: 36 x 4 [1Q]
# Key:      .model [3]
   .model Quarter      Gas .mean
   <chr>   <qtr>      <dist> <dbl>
1 drift  2007 Q3 t(N(7.4, 0.077)) 240.
2 drift  2007 Q4 t(N(7.4, 0.15)) 247.
3 drift  2008 Q1 t(N(7.4, 0.23)) 253.
4 drift  2008 Q2 t(N(7.4, 0.31)) 260.
5 drift  2008 Q3 t(N(7.5, 0.39)) 267.
6 drift  2008 Q4 t(N(7.5, 0.47)) 274.
7 drift  2009 Q1 t(N(7.5, 0.55)) 281.
8 drift  2009 Q2 t(N(7.6, 0.64)) 288.
9 drift  2009 Q3 t(N(7.6, 0.72)) 296.
```



```
10 drift 2009 Q4 t(N(7.6, 0.8)) 303.
```

```
# i 26 more rows
```

```
gas_fc |>
  autoplot(aus_prod_recent) +
  facet_wrap(vars(.model), scale = "free_y", ncol = 1)
```



El modelo seasonal naïve copia y pega el último año, su pronóstico es el igual al último año y no capturó la tendencia de la serie.

$$MAE = mean(|e_t|)$$

Estas medidas no pueden usarse para comparar pronósticos entre distintas series por la escala de los datos. Sin embargo, se puede usar el error porcentual como el MAPE, el error absoluto medio porcentual:

$$MAPE = mean(|P_t|)$$

```
gas_fc |>
  accuracy(aus_production) |>
  arrange(RMSE)
```

A tibble: 3 × 10


| | .model | .type | ME | RMSE | MAE | MPE | MAPE | MASE | RMSSE | ACF1 |
|---|--------|-------|-------|-------|-------|-------|-------|-------|-------|---------|
| | <chr> | <chr> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 | snaive | Test | 8.12 | 10.6 | 9.42 | 3.72 | 4.28 | 1.70 | 1.39 | -0.0968 |
| 2 | drift | Test | -56.5 | 64.7 | 57.3 | -26.5 | 26.8 | 10.4 | 8.51 | 0.357 |
| 3 | media | Test | 127. | 128. | 127. | 56.7 | 56.7 | 22.9 | 16.9 | -0.0534 |

Pronóstico por descomposición

```
gas_train |>
  model(
    stl = STL(box_cox(Gas, gas_lambda) ~ season(window= "periodic"), robust = TRUE)
  ) |> #robust = true es para que los outliers el efecto se vaya al componente residual
  components() |>
  autoplot()
```

STL decomposition

`box_cox(Gas, gas_lambda)` = trend + season_year + remainder



```
gas_dcmp <- gas_train |>
  model(
    dcmp = decomposition_model(
      STL(box_cox(Gas, gas_lambda) ~ season(window = "periodic"), robust = TRUE),
      RW(season_adjust ~ drift()),
      SNAIVE(season_year)
```

①

②

③

④

```
)
)

gas_dcmp
```

- ① `decomposition_model()` define que se realizará un pronóstico a partir de una descomposición.
- ② Primero se define cómo se realizará la descomposición. En este caso, con STL, con ajuste robusto y con componente estacional periódica.
- ③ Luego se define el modelo para la serie desestacionalizada, `season_adjust`.
- ④ Finalmente, se define el modelo para la serie estacional, `season_year`. Si este componente no se especifica, R va a utilizar `SNAIVE()` por default.

```
# A mable: 1 x 1
```

```
      dcmp
      <model>
```

```
1 <STL decomposition model>
```



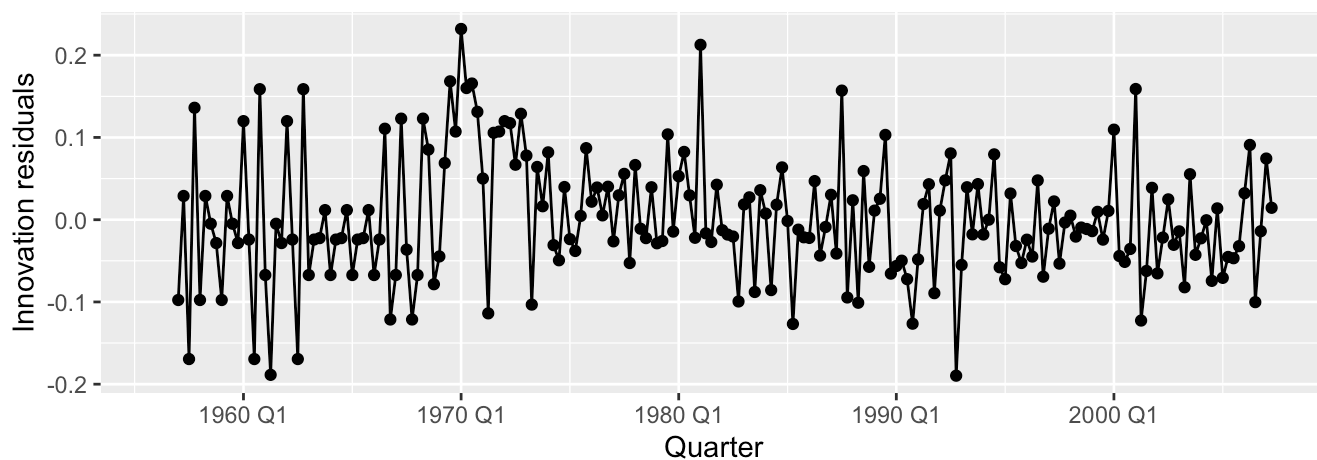
$H_0 : Q = 0$ ► No autocorrelación

$H_i : Q \neq 0$ ► Sí hay autocorrelación

Rechazar o no una hipótesis H_0 depende de una α , donde α es la probabilidad de cometer un error de tipo I (Probabilidad de rechazar H_0 cuando es verdadera). El error de tipo II es No rechazar la H_0 cuando es falsa.

si p-value < α rechazo H_0

```
gas_dcmp |>
  gg_tsresiduals()
```



```
gas_dcmp |>
  augment() |>
  features(.innov, ljung_box, lag = 8)
```

```
# A tibble: 1 × 3
  .model lb_stat lb_pvalue
  <chr>   <dbl>   <dbl>
1 dcmp    33.6 0.0000473
```

```
gas_dcmp |>
  accuracy()
```


```
# A tibble: 1 × 10
  .model .type      ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
  <chr>  <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 dcmp  Training -0.381  4.28  2.83 -0.108  4.02  0.512  0.563 -0.184
```

```
gas_dcmp_fc <- gas_dcmp |>
  forecast(h = "3 years")

gas_dcmp_fc
```

```
# A fable: 12 x 4 [1Q]
# Key:      .model [1]
. model Quarter      Gas .mean
<chr>    <qtr>      <dist> <dbl>
1 dcmp   2007 Q3  t(N(7.5, 0.0058)) 258.
2 dcmp   2007 Q4  t(N(7.2, 0.012)) 215.
3 dcmp   2008 Q1  t(N(7.1, 0.018)) 201.
4 dcmp   2008 Q2  t(N(7.4, 0.024)) 249.
5 dcmp   2008 Q3  t(N(7.6, 0.03)) 274.
6 dcmp   2008 Q4  t(N(7.3, 0.036)) 228.
7 dcmp   2009 Q1  t(N(7.2, 0.042)) 213.
8 dcmp   2009 Q2  t(N(7.5, 0.048)) 264.
9 dcmp   2009 Q3  t(N(7.7, 0.054)) 291.
10 dcmp  2009 Q4  t(N(7.4, 0.061)) 243.
11 dcmp  2010 Q1  t(N(7.3, 0.067)) 227.
12 dcmp  2010 Q2  t(N(7.6, 0.074)) 281.
```

```
gas_dcmp_fc |>
  autoplot(aus_production |> filter_index("1999 Q1" ~ .))
```



```
gas_fc_full <- gas_fc |>
  full_join(gas_dcmp_fc)
```

Joining with `by = join_by(.model, Quarter, Gas, .mean)`


```
gas_fc_full
```

```
# A tibble: 48 x 4 [10]
# Key:   .model [4]
  .model Quarter      Gas .mean
  <chr>   <qtr>      <dist> <dbl>
1 dcmp   2007 Q3  t(N(7.5, 0.0058)) 258.
2 dcmp   2007 Q4  t(N(7.2, 0.012)) 215.
3 dcmp   2008 Q1  t(N(7.1, 0.018)) 201.
4 dcmp   2008 Q2  t(N(7.4, 0.024)) 249.
5 dcmp   2008 Q3  t(N(7.6, 0.03)) 274.
6 dcmp   2008 Q4  t(N(7.3, 0.036)) 228.
7 dcmp   2009 Q1  t(N(7.2, 0.042)) 213.
8 dcmp   2009 Q2  t(N(7.5, 0.048)) 264.
9 dcmp   2009 Q3  t(N(7.7, 0.054)) 291.
10 dcmp  2009 Q4  t(N(7.4, 0.061)) 243.
# i 38 more rows
```

```
gas_fc_full |>
  accuracy(aus_production) |>
  arrange(RMSE)
```

```
# A tibble: 4 x 10
  .model .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
  <chr>  <chr>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 snaive Test    8.12  10.6   9.42  3.72  4.28  1.70  1.39 -0.0968
2 dcmp   Test   -23.5  25.9  23.5 -10.4  10.4  4.24  3.41  0.389
3 drift  Test  -56.5  64.7  57.3 -26.5  26.8  10.4  8.51  0.357
4 media  Test  127.  128.  127.  56.7  56.7  22.9  16.9 -0.0534
```

```
gas_fc_full |>
  autoplot(aus_production |> filter_index("1999 Q1" ~ .), level = NULL)
```



```
gas_fit_full <- gas_fit |>
  cross_join(gas_dcmp) |>
  mutate(combinado = (snaive + dcmp)/2)


gas_fit_full
```

A mable: 1 x 5


| | drift | snaive | media | dcmp | combinado |
|---|---------------|----------|---------|---------------------------|---------------|
| | <model> | <model> | <model> | <model> | <model> |
| 1 | <RW w/ drift> | <SNAIVE> | <MEAN> | <STL decomposition model> | <COMBINATION> |

```
gas_fcst_full <- gas_fit_full |>
  forecast(h = "3 years")

gas_fcst_full |>
  autoplot(aus_production |> filter_index("1999 Q1" ~ .), level = NULL)
```



```
gas_fcst_full |>
  filter(.model == "combinado") |>
  autoplot(aus_production |> filter_index("1999 Q1" ~ .))
```

```
gas_fcst_full |>
  accuracy(aus_production) |>
  arrange(RMSE)
```

A tibble: 5 × 10

| .model | .type | ME | RMSE | MAE | MPE | MAPE | MASE | RMSSE | ACF1 |
|-------------|-------|-------|-------|-------|-------|-------|-------|-------|---------|
| <chr> | <chr> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 combinado | Test | -7.67 | 10.1 | 7.86 | -3.36 | 3.46 | 1.42 | 1.33 | 0.0294 |
| 2 snaive | Test | 8.12 | 10.6 | 9.42 | 3.72 | 4.28 | 1.70 | 1.39 | -0.0968 |
| 3 dcmp | Test | -23.5 | 25.9 | 23.5 | -10.4 | 10.4 | 4.24 | 3.41 | 0.389 |
| 4 drift | Test | -56.5 | 64.7 | 57.3 | -26.5 | 26.8 | 10.4 | 8.51 | 0.357 |
| 5 media | Test | 127. | 128. | 127. | 56.7 | 56.7 | 22.9 | 16.9 | -0.0534 |

ETS

```
ses <- gas_train |>
  model(
    ses = ETS(Gas ~ error("A") + trend("N") + season("N"))
  )

ses
```

```
# A mable: 1 x 1
```

```
ses
```

```
<model>
```

```
1 <ETS(A,N,N)>
```

1. `ETS()` es la función para estimar modelos de suavización exponencial y se deben definir 3 argumentos: error, tendencia y estacionalidad.
2. Tenemos dos opciones: "A" para error aditivo, y "M" para multiplicativo
3. Para especificar que no queremos ni tendencia, ni estacionalidad, ponemos "N" en ambos casos.

```
report(ses)
```

```
Series: Gas
```

```
Model: ETS(A,N,N)
```

```
Smoothing parameters:
```

```
alpha = 0.2434761
```

```
Initial states:
```

```
l[0]
```

```
5.9707
```

```
sigma^2: 261.8623
```


```
AIC      AICc      BIC
```

```
2248.503 2248.622 2258.487
```

```
ses |>
```

```
  forecast(h = "3 years") |>
```

```
  autoplot(aus_prod_recent)
```



Tendencia lineal de HOLT

```
holt <- gas_train |>
  model(
    holt = ETS(Gas ~ error("A") + trend("A") + season("N"))
  )

holt
```

```
# A mable: 1 x 1
      holt
  <model>
1 <ETS(A,A,N)>
```

```
report(holt)
```

```
Series: Gas
Model: ETS(A,A,N)
Smoothing parameters:
  alpha = 0.1117699
  beta  = 0.01133185
```


Initial states:

$l[0]$ $b[0]$
6.681936 0.01735882

σ^2 : 236.0876

AIC AICc BIC
2229.129 2229.429 2245.768

```
holt |>
  forecast(h = "3 years") |>
  autoplot(aus_prod_recent)
```



Tendencia amortiguada

Holt winters

```
hw <- gas_train |>
  model(
    hw = ETS(Gas ~ error("M") + trend("Ad", phi = 0.8) + season("M")) #poner error = esta
```

)

```
report(hw)
```

Series: Gas

Model: ETS(M,Ad,M)

Smoothing parameters:

alpha = 0.5788133

beta = 0.2907041

gamma = 0.0736501

phi = 0.8

Initial states:

| l[0] | b[0] | s[0] | s[-1] | s[-2] | s[-3] |
|----------|------------|-----------|----------|----------|-----------|
| 5.837584 | 0.08005418 | 0.9305848 | 1.182244 | 1.069461 | 0.8177104 |


sigma^2: 0.0037

| AIC | AICc | BIC |
|----------|----------|----------|
| 1565.797 | 1566.715 | 1595.748 |

```
hw |>
```

```
forecast(h = "3 years") |>
```

```
autoplot(aus_prod_recent)
```



Comparando modelos

```
STLF <- decomposition_model(
  STL(box_cox(Gas, gas_lambda) ~ season(window = "periodic"), robust=TRUE),
  RW(season_adjust ~ drift())
)

STLF_ets <- decomposition_model(
  STL(box_cox(Gas, gas_lambda) ~ season(window = "periodic"), robust=TRUE),
  ETS(season_adjust ~ error("A") + trend("Ad") + season("N"))
)

gas_fit_todos <- gas_train |>
  model(
    hw = ETS(Gas ~ error("M") + trend("Ad", phi = 0.8) + season("M")),
    snaive = SNAIVE(box_cox(Gas, gas_lambda)),
    hw_boxcox = ETS(box_cox(Gas, lambda = gas_lambda) ~ error("A") + trend("Ad") + season("M")),
    stlf = STLF,
    stlf_ets = STLF_ets
  )

gas_fit_todos
```

```
# A mable: 1 x 5
      hw      snaive      hw_boxcox      stlf
  <model> <model>      <model>      <model>
1 <ETS(M,Ad,M)> <SNAIVE> <ETS(A,Ad,A)> <STL decomposition model>
# i 1 more variable: stlf_ets <model>
```

```
gas_fc_todos <- gas_fit_todos |>
  forecast(h = "3 years")

gas_fc_todos |>
  accuracy(aus_production) |>
  arrange(RMSE)
```

```
# A tibble: 5 x 10
  .model .type ME RMSE MAE MPE MAPE MASE RMSSE ACF1
  <chr>   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 snaive Test  8.12 10.6  9.42  3.72  4.28  1.70  1.39 -0.0968
2 hw      Test -14.3 15.5 14.3 -6.32  6.32  2.58  2.04  0.0430
3 stlf_ets Test -16.2 18.0 16.2 -7.19  7.19  2.93  2.37  0.201
4 hw_boxcox Test -16.6 18.3 16.6 -7.37  7.37  3.00  2.40  0.231
5 stlf     Test -23.5 25.9 23.5 -10.4 10.4  4.24  3.41  0.389
```

```
p <- gas_fc_todos |>
  autoplot(aus_prod_recent, level = NULL)
```

```
plotly::ggplotly(p)
```

